

CS344: OPERATING SYSTEMS LAB

Assignment 3

Group 13

Adarsh Kumar, 190101002

Tanishq Katare, 190101093

Upender Dahiya, 190101096

Keshav Chourasiya, 190101045

PART A

To implement the lazy memory allocation we need to change the following files:

- **sysproc.c** : (File already provided) The function **sbrk()** is used to allocate the physical memory. So, for lazy memory allocation implementation we just need to comment the **growproc()** function calling part which is responsible for this, and also keep track of the size by incrementing the **sz** attribute of the process

```
45  int
46  sys_sbrk(void)
47  {
48      int addr;
49      int n;
50
51      if(argint(0, &n) < 0)
52          return -1;
53      addr = myproc()->sz;
54      myproc()->sz += n;
55      // if(growproc(n) < 0)
56      //     return -1;
57      return addr;
58  }
```

- **trap.c** : An additional condition was added in the switch-case section. A page fault occurs when **tf->trapno** equals **T_PGFLT**, so therefore in this case **kalloc()** and **mappages()** were called to allocate the required space or to inform that no more memory is available (code reused from **allocvm()** in **vm.c**)

```
82  case T_PGFLT:{
83      char *mem;
84      mem = kalloc();
85      if (mem == 0) {
86          cprintf("Out of Memory.\n");
87      } else {
88          memset(mem, 0, PGSIZE);
89          // creating page table entry
90          uint a = PGROUNDOWN(rcr2());
91          mappages(myproc()->pgdir, (char *)a, PGSIZE, V2P(mem), PTE_W | PTE_U);
92      }
93      break;
94  }
```

- **vm.c** : Return type of function **mappages()** was changed from **static int** to **int** so that it can be used in **trap.c**

PART B

ANSWERING TO QUESTIONS

- How does the kernel know which physical pages are used and unused?

Ans: xv6 maintains a linked list of free pages in **kalloc.c** called **kmem**. This list is empty first and xv6 initialises this list with 4MB free pages from **main.c** file where it calls **kinit** for kernel initialisation

- What data structures are used to answer this question?

Ans: xv6 uses **singly linked list** as the data structure for these free pages

- Where do these reside?

Ans: This linked list is declared inside **kalloc.c** inside a structure **kmem**. Every node is of the type **struct run** which is also defined inside **kalloc.c**

- Does xv6 memory mechanism limit the number of user processes. If so, what is the lowest number of processes xv6 can 'have' at the same time (assuming the kernel requires no memory whatsoever)?

Ans: Under file **param.h** the macro **NPROC** set as **64**, limits the number of user processes that can be present simultaneously in the ptable. Lowest number of process xv6 can have is **1**, there cannot be 0 processes after boot since all user interactions need to be done using user processes which are forked from **initproc** and **sh**

TASK 1

In *proc.c*, the *create_kernel_process()* function was created and this kernel process will stay in kernel mode throughout its lifetime. We don't need to initialise its trapframe, user space, or the user sector of its page table.

The address of the next instruction is stored in the process context's *eip* register. We want the process to begin at the point of entrance (which is a function pointer given as an argument) therefore we set the context's *eip* value to the entry point, note entry point is the address of a function. *allocproc* assigns a position in ptable to the process. The kernel section of the process page table, which translates virtual addresses above *KERNBASE* to physical locations between 0 and *PHYSTOP*, is set up using *setupkvm*.

```
441 void create_kernel_process(const char *name, void (*entrypoint)()){
442     struct proc *p = allocproc();
443     if(p == 0)
444         panic("create_kernel_process failed");
445
446     //Setting up kernel page table using setupkvm
447     if((p->pgdir = setupkvm()) == 0)
448         panic("setupkvm failed");
449     p->context->eip = (uint)entrypoint;
450
451     safestrcpy(p->name, name, sizeof(p->name));
452
453     acquire(&ptable.lock);
454     p->state = RUNNABLE;
455     release(&ptable.lock);
456 }
```

TASK 2

Dividing this task into modules and therefore firstly, we required a process queue to keep track of the processes that were denied more memory due to lack of free pages. For this we generated the ***rq*** circular queue struct and the queue ***rqueue*** to contain the processes which have swap out requests. We've also created the ***rpush()*** and ***rpop()*** functions.

Code snippets from ***proc.c***:

```
155 struct rq{
156     struct spinlock lock;
157     struct proc* queue[NPROC];
158     int s;
159     int e;
160 };
```

```
178 int rpush(struct proc *p){
179     acquire(&rqueue.lock);
180     if((rqueue.e+1)%NPROC==rqueue.s){
181         release(&rqueue.lock);
182         return 0;
183     }
184     rqueue.queue[rqueue.e]=p;
185     rqueue.e++;
186     (rqueue.e)%=NPROC;
187     release(&rqueue.lock);
188     return 1;
189 }
```

```
165 struct proc* rpop(){
166     acquire(&rqueue.lock);
167     if(rqueue.s==rqueue.e){
168         release(&rqueue.lock);
169         return 0;
170     }
171     struct proc *p=rqueue.queue[rqueue.s];
172     rqueue.s+=1;
173     rqueue.s%=NPROC;
174     release(&rqueue.lock);
175     return p;
176 }
```

We will access this queue using a lock which we have initialised in ***pinit***. In ***userinit***, we've set the starting values of ***s*** and ***e*** to zero. We introduced prototypes in ***defs.h*** as well, because the queue and functions related to it are used in other files.

```
460 void
461 userinit(void)
462 {
463     acquire(&rqueue.lock);
464     rqueue.s=0;
465     rqueue.e=0;
466     release(&rqueue.lock);
467
468     acquire(&rqueue2.lock);
469     rqueue2.s=0;
470     rqueue2.e=0;
471     release(&rqueue2.lock);
```

```
343 void
344 pinit(void)
345 {
346     initlock(&ptable.lock, "ptable");
347     initlock(&rqueue.lock, "rqueue");
348     initlock(&sleeping_channel_lock, "sleeping_channel");
349     initlock(&rqueue2.lock, "rqueue2");
350 }
```

If ***kalloc*** is unable to allocate pages to a process, it now returns 0. This informs ***allocuvmm*** that the requested memory(mem=0) wasn't allocated. To begin, we must set the process state to 'sleeping'. The process sleeps on a unique sleeping channel called ***sleeping_channel*** which is protected by a lock ***sleeping_channel_lock***. When the system boots, ***sleeping_channel_count*** is used for the edge cases. Next, we must add the current process to the ***rqueue***.

Code snippet of allocvm function:

```
244
245 //SLEEP
246 myproc()->state=SLEEPING;
247 acquire(&sleeping_channel_lock);
248 myproc()->chan=sleeping_channel;
249 sleeping_channel_count++;
250 release(&sleeping_channel_lock);
251
252 rpush(myproc());
253 if(!swap_out_process_exists){
254     swap_out_process_exists=1;
255     create_kernel_process("swap_out_process", &swap_out_process_function);
256 }
257
```

Code snippet from kalloc.c:

create_kernel_process here creates a swapping out kernel process to allocate a page for this process if it doesn't already exist. The **swap_out_process_exists** variable is set to 0 when the swap out process is completed. This is done to avoid the creation of multiple swap out operations.

Next, we develop a mechanism that wakes up all processes sleeping on **sleeping_channel** whenever free pages become available. Now we make the following changes to **kfree** function in **kalloc.c**. Basically, all processes that were preempted due to lack of availability of pages were sent sleeping on the sleeping channel. The **wakeup()** system call is used to wake up any processes that are currently sleeping on the sleeping channel.

```
60 void
61 kfree(char *v)
62 {
63     struct run *r;
64
65     if((uint)v % PGSIZE || v < end || V2P(v) >= PHYSTOP)
66         panic("kfree");
67
68     // Fill with junk to catch dangling refs.
69     // memset(v, 1, PGSIZE);
70     for(int i=0;i<PGSIZE;i++){
71         v[i]=1;
72     }
73
74     if(kmem.use_lock)
75         acquire(&kmem.lock);
76     r = (struct run*)v;
77     r->next = kmem.freelist;
78     kmem.freelist = r;
79     if(kmem.use_lock)
80         release(&kmem.lock);
81
82     //Wake up processes sleeping on sleeping channel.
83     if(kmem.use_lock)
84         acquire(&sleeping_channel_lock);
85     if(sleeping_channel_count){
86         wakeup(sleeping_channel);
87         sleeping_channel_count=0;
88     }
89     if(kmem.use_lock)
90         release(&sleeping_channel_lock);
91
92 }
```

Now, we will explain the **swapping out process**. The entry point for the swapping out process is in the **swap_out_process_function**.

The process runs a loop until the swap out requests queue (**rqueue1**) is non empty. When the queue is empty, a set of instructions are executed for the termination of **swap_out_process**.

```

220 void swap_out_process_function(){
221     acquire(&rqueue.lock);
222     while(rqueue.s!=rqueue.e){
223         struct proc *p=rpop();
224         pde_t* pd = p->pgdir;
225         for(int i=0;i<NPENTRIES;i++){
226             //skip page table if accessed. chances are high, not every pa
227             if(pd[i]&PTE_A)
228                 continue;
229             pte_t *pgtab = (pte_t*)P2V(PTE_ADDR(pd[i]));
230             for(int j=0;j<NPENTRIES;j++){
231                 //Skip if found
232                 if((pgtab[j]&PTE_A) || !(pgtab[j]&PTE_P))
233                     continue;
234                 pte_t *pte=(pte_t*)P2V(PTE_ADDR(pgtab[j]));
235
236                 //for file name
237                 int pid=p->pid;
238                 int virt = ((1<<22)*i)+((1<<12)*j);
239
240                 //file name
241                 char c[50];
242                 int_to_string(pid,c);
243                 int x=strlen(c);
244                 c[x]='_';
245                 int_to_string(virt,c+x+1);
246                 safestrcpy(c+strlen(c),".swp",5);
247
248                 // file management
249                 int fd=proc_open(c, 0_CREATE | 0_RDWR);
250                 if(fd<0){
251                     cprintf("error creating or opening file: %s\n", c);
252                     panic("swap_out_process");
253                 }

```

```

248 // file management
249 int fd=proc_open(c, 0_CREATE | 0_RDWR);
250 if(fd<0){
251     cprintf("error creating or opening file: %s\n", c);
252     panic("swap_out_process");
253 }
254 if(proc_write(fd,(char *)pte, PGSIZE) != PGSIZE){
255     cprintf("error writing to file: %s\n", c);
256     panic("swap_out_process");
257 }
258 proc_close(fd);
259 kfree((char*)pte);
260 memset(&pgtab[j],0,sizeof(pgtab[j]));
261 //mark this page as being swapped out.
262 pgtab[j]=(pgtab[j])^(0x080);
263 break;
264 }
265 }
266 }
267 release(&rqueue.lock);
268 struct proc *p;
269 if((p=myproc())==0)
270     panic("swap out process");
271
272 swap_out_process_exists=0;
273 p->parent = 0;
274 p->name[0] = '*';
275 p->killed = 0;
276 p->state = UNUSED;
277 sched();
278 }

```

The loop begins by removing the first process from **rqueue** and determining a victim page in the page table using the **LRU policy**. We derive the physical address for each secondary page table by iterating through each entry in the process page table, **pgdir**. We go through the page table for each secondary page table, looking at the accessed bit (A) on each entry. When the scheduler switches the process's context, all accessed bits are cleared. Since we are doing this, the accessed bit seen by **swap_out_process_function** will indicate whether the entry was accessed in the last iteration of the process. This code resides in the scheduler and it basically unsets every accessed bit in the process page table and its secondary page tables.

Now, back to the **swap_out_process_function**. As soon as the function finds a secondary page table entry with the accessed bit unset, it chooses this entry's physical page number as the victim page. After that, this page is then swapped out and saved on the drive. To name the file that saves this page, we use the process **pid**(line 237) and **virtual address** of the page(line 238) to be erased. We need to write the contents of the victim page to the file with the name **<pid>_<virt>.swp**.

But we have a problem here. The filename is saved in a string named **c**. Calls to the file system cannot be made from **proc.c**. So we copied the **open**, **write**, **read**, **close** functions from **sysfile.c** to **proc.c**, and updated and renamed them to **proc_open**, **proc_read**, **proc_write**, **proc_close**, and so on so we could use them in **proc.c**. Here are some examples:

```
20 | int proc_close(int fd){
21 |     struct file *f;
22 |     if(fd < 0 || fd >= NOFILE || (f=myproc()->ofile[fd]) == 0)
23 |         return -1;
24 |     myproc()->ofile[fd] = 0;
25 |     fileclose(f);
26 |     return 0;
27 | }
28 |
29 | int proc_write(int fd, char *p, int n){
30 |     struct file *f;
31 |     if(fd < 0 || fd >= NOFILE || (f=myproc()->ofile[fd]) == 0)
32 |         return -1;
33 |     return filewrite(f, p, n);
34 | }
```

Now, using these functions, we write back a page to storage. We open a file (using **proc_open**) with **O_CREATE** and **O_RDWR** permissions (we have imported **fcntl.h** with these macros). **O_CREATE** creates this file if it doesn't exist and **O_RDWR** refers to read/write. The file descriptor is stored in an integer called **fd**. Using this file descriptor, we write the page to this file using **proc_write**. The page is then added to the free page queue using **kfree**, making it available for usage (note that when **kfree** adds a page to the free queue, it also wakes up all processes sleeping on the sleeping channel). We then use **memset** to clear the page table entry as well. For Task 3, we need to know if the page that caused a fault was swapped out or not. For marking the page as swapped out, we set the 8th bit from the right (2^7) in the secondary page table entry using **xor**.

TASK 3

To begin, we must first make a swap in the request queue. In **proc.c**, we created a swap in request queue called **rqueue2** using the same **struct rq**. In **defs.h**, we additionally specify an extern prototype for **rqueue2**. We developed the matching functions for **rqueue2** for push and pop operations named **rpop2()** and **rpush2()** in **proc.c**. In **pinit**, we also initialised its lock and in **userinit**, we additionally set the **s** and **e** attributes.

Then, in **proc.h**, we added an extra entry named **addr** to the **struct proc**. This item will inform the swapping in function about the virtual address at which page fault was found:

```
50 struct proc {
51     char name[16];           // Process name (debugging)
52     int addr;                // Virtual address of pagefault
53 };
```

Next, we need to handle page fault (**T_PGFLT**) traps raised in **trap.c**. We do it in a function called **handlePageFault()**:

```
104 case T_PGFLT:
105     handlePageFault();
106 break;
```

```
19 void handlePageFault(){
20     int addr=rcr2();
21     struct proc *p=myproc();
22     acquire(&swap_in_lock);
23     sleep(p,&swap_in_lock);
24     pde_t *pde = &(p->pgdir)[PDX(addr)];
25     pte_t *pgtab = (pte_t*)P2V(PTE_ADDR(*pde));
26
27     if((pgtab[PTX(addr)]&0x080)){
28         //This means that the page was swapped out.
29         //virtual address for page
30         p->addr = addr;
31         rpush2(p);
32         if(!swap_in_process_exists){
33             swap_in_process_exists=1;
34             create_kernel_process("swap_in_process", &swap_in_process_function);
35         }
36     } else {
37         exit();
38     }
39 }
```

In **handlePageFault()**, we use **rcr2()** to identify the virtual address where the page fault occurred, exactly as we did in Part A. We then use a new lock called **swap_in_lock** to put the current process to sleep. Then we get the page table entry analogous to this address. Now we have to find if this page was swapped out or not.

When we swapped out a page in Task 2, we set its page table entry's bit of 7th order (2^7).

To detect if the page was swapped out or not, we apply bitwise & with 0x080 to check its 7th order bit. If it's set, **swap_in_process** is initiated. Otherwise, we may safely suspend the process using **exit()**.

Now we'll go over the swapping in procedure. As you can see in `handlePageFault`, the

`swap_in_process_function` is the entry point for the swapping out process.

Until **`rqueue2`** is not empty, the file management procedure loops. In the loop, it gets the file name by popping a process from the queue and extracting its **`pid`** and **`addr`** values. It then uses **`int_to_string`** to construct the filename in a string named "c". Then, given file descriptor **`fd`**, it used **`proc_open`** to open this file in read-only mode. Using **`kalloc`**, we then allocate a free frame to this process. Using **`proc_read`**, we read from the file with the `fd`, file descriptor into this free frame. By deleting the **`static`** keyword from **`vm.c`** and declaring a prototype in **`proc.c`**, we make **`mappages`** accessible to **`proc.c`**. The page corresponding to **`addr`** is then mapped with the physical page obtained using **`kalloc`** and read into using **`mappages`**. Then we wake up, which is a procedure in which we allocate a new page to correct the page fault. We perform the kernel process termination instructions once the loop is complete.

```
287 void swap_in_process_function(){
288     acquire(&rqueue2.lock);
289     while(rqueue2.s!=rqueue2.e){
290         struct proc *p=rpop2();
291         int pid=p->pid;
292         int virt=PTE_ADDR(p->addr);
293
294         char c[50];
295         int_to_string(pid,c);
296         int x=strlen(c);
297         c[x]='_';
298         int_to_string(virt,c+x+1);
299         safestrcpy(c+strlen(c),".swp",5);
300
301         int fd=proc_open(c,0_RDONLY);
302         if(fd<0){
303             release(&rqueue2.lock);
304             cprintf("could not find page file in memory: %s\n", c);
305             panic("swap_in_process");
306         }
307         char *mem=kalloc();
308         proc_read(fd,PGSIZE,mem);
309
310         if(mappages(p->pgdir, (void *)virt, PGSIZE, V2P(mem), PTE_W|PTE_U)<0){
311             release(&rqueue2.lock);
312             panic("mappages");
313         }
314         wakeup(p);
315     }
316
317     release(&rqueue2.lock);
318     struct proc *p;
319     if((p=myproc())==0)
320         panic("swap_in_process");
321
322     swap_in_process_exists=0;
323     p->parent = 0;
324     p->name[0] = '*';
325     p->killed = 0;
326     p->state = UNUSED;
327     sched();
328 }
```

TASK 4

- In this task, a test user-program is created named **memtest**.
- This program has main process that forks and creates **20 child processes**
- Each of this child process is looping through **10 iterations** and at each iteration, **4096B** of memory is being allocated using the **malloc()** function
- As asked a random mathematical expression is used at every ith iteration to store some value at ith index of the array
- After filling the values in the memory, to validate used the same mathematical function to compare the stored values
- Since this user program is a new one hence under **UPROGS** and **EXTRA** in the **Makefile** we need to update the appropriate function name

We can notice from the output that the sanity test gives the correct output since the unmatched column that denotes the total unmatched bytes are all 0B and hence shows that the indices store the correct value

Now upon testing on different values of **PHYSTOP** defines under the file **memlayout.h** The correct implementation should show identical results.

Therefore on changing the default value of PHYSTOP i.e. **0xe000000(224MB)** to something smaller such as **0x0400000(4MB)** the output doesn't change indicating that the implementation is correct

```
memtest      2 18 16972
console      3 19 0
$ memtest
```

	Iteration	Matched	Unmatched
Child 1	1	4096B	0B
Child 1	2	4096B	0B
Child 1	3	4096B	0B
Child 1	4	4096B	0B
Child 1	5	4096B	0B
Child 1	6	4096B	0B
Child 1	7	4096B	0B
Child 1	8	4096B	0B
Child 1	9	4096B	0B
Child 1	10	4096B	0B
Child 2	1	4096B	0B
Child 2	2	4096B	0B
Child 2	3	4096B	0B
Child 2	4	4096B	0B
Child 2	5	4096B	0B
Child 2	6	4096B	0B
Child 2	7	4096B	0B
Child 2	8	4096B	0B
Child 2	9	4096B	0B
Child 2	10	4096B	0B
Child 3	1	4096B	0B
Child 3	2	4096B	0B
Child 3	3	4096B	0B
Child 3	4	4096B	0B
Child 3	5	4096B	0B
Child 3	6	4096B	0B
Child 3	7	4096B	0B
Child 3	8	4096B	0B
Child 3	9	4096B	0B
Child 3	10	4096B	0B
Child 4	1	4096B	0B
Child 4	2	4096B	0B
Child 4	3	4096B	0B
Child 4	4	4096B	0B
Child 4	5	4096B	0B